

# Machine Learning (CE 40477)

Fall 2024

Ali Sharifi-Zarchi

CE Department  
Sharif University of Technology

September 30, 2024



- 1 Generalization
- 2 Probabilistic regression
- 3 References



## Generalization Overview

**Main Idea:** The ability of a model to perform well on unseen data

- **Training Set:**  $D = \{(x_i, y_i)\}_{i=1}^n$
- **Test Set:** New data not seen during training
- **Cost Function:** Measures how well the model fits data

$$J(w) = \sum_{i=1}^n (y^{(i)} - h_w(x^{(i)}))^2$$

- **Objective:** Minimize the cost function on unseen data (generalization error)

# Expected Test Error

**Definition:** Expected performance on unseen data

- Test data sampled from the same distribution  $p(x, y)$

$$J(w) = \mathbb{E}_{p(x,y)} [(y - h_w(x))^2]$$

- Approximate using test set  $\hat{J}(w)$
- Generalization error is the gap between training and test performance.

# Training vs Test Error

**Key Concept:** Training error measures fit on known data, test error on unseen data

- **Training (empirical) error:**

$$J_{\text{train}}(w) = \frac{1}{n} \sum_{i=1}^n \left( y^{(i)} - h_w(x^{(i)}) \right)^2$$


- **Test error:**

$$J_{\text{test}}(w) = \frac{1}{m} \sum_{i=1}^m \left( y_{\text{test}}^{(i)} - h_w(x_{\text{test}}^{(i)}) \right)^2$$

- **Goal:** Minimize the test error (generalization).

# Overfitting Definition


**Concept:** A model fits the training data well but performs poorly on the test set


$$J_{\text{train}}(w) \ll J_{\text{test}}(w)$$

- Causes: Model too complex, **high variance**
- Consequence: Captures noise in training data, fails on unseen data

# Underfitting Definition


**Concept:** The model is too simple and cannot capture the structure of the data

  $J_{\text{train}}(w) \approx J_{\text{test}}(w) \gg 0$


- Causes: Model lacks complexity, **high bias**
- Consequence: Poor fit on both training and test data



# Generalization: polynomial regression




Degree of 1




Degree of 3

Example adapted from slides of Dr. Soleymani, ML course, Sharif University of technology.

# Overfitting: polynomial regression





Degree of 5





Degree of 7

# Polynomial regression with various degrees: example




Figures adapted from Machine Learning and Pattern Recognition, Bishop

# Polynomial regression with various degrees: example (cont.)



Figures adapted from Machine Learning and Pattern Recognition, Bishop

# Root mean squared error



$$E_{RMS} = \sqrt{\frac{\sum_{i=1}^n (y^{(i)} - f(\mathbf{x}^{(i)}; \mathbf{w}))^2}{n}}$$

# Bias-Variance Decomposition

## Generalization error decomposition:

$$\mathbb{E}[(y - h_w(x))^2] = (\text{Bias})^2 + \text{Variance} + \text{Noise}$$

- **Bias:** Error due to simplifying assumptions in the model

$$\text{Bias}(x) = \mathbb{E}[h_w(x)] - f(x)$$

- **Variance:** Sensitivity of the model to training data

$$\text{Variance}(x) = \mathbb{E}[(h_w(x) - \mathbb{E}[h_w(x)])^2]$$

- **Noise:** Irreducible error from the inherent randomness in data

# Bias-Variance Decomposition Proof

Assume  $f(x)$  is the ground truth and observation  $y$  is a noisy observation  $y = f(x) + \epsilon$  where  $\epsilon \sim \mathcal{N}(0, \sigma^2)$ . We start with the definition of the expected squared error, which is:

$$\begin{aligned}\mathbb{E}_{data} \left[ (\hat{f}(x) - y)^2 \right] &= \mathbb{E}_{data} \left[ (\hat{f}(x) - f(x) + \epsilon)^2 \right] \\ &= \mathbb{E} \left[ (\hat{f}(x) - f(x))^2 - 2\epsilon (\hat{f}(x) - f(x)) + \epsilon^2 \right]\end{aligned}$$

Since we assume the noise  $\epsilon$  has zero mean and variance  $\sigma^2$ , the term  $\mathbb{E}[\epsilon] = 0$ , and thus:

$$\mathbb{E}[\epsilon^2] = \sigma^2$$

Since  $\mathbb{E}[\epsilon] = 0$  and  $\epsilon$  is independent of the parenthesis, we can write:

$$\mathbb{E} \left[ -2\epsilon (\hat{f}(x) - f(x)) \right] = 0$$

# Bias-Variance Decomposition Proof (cont.)

Now, we decompose the squared difference  $(\hat{f}(x) - f(x))^2$  as follows:

$$\mathbb{E} \left[ (\hat{f}(x) - f(x))^2 \right] = \mathbb{E} \left[ (\hat{f}(x) - \mathbb{E} [\hat{f}(x)] + \mathbb{E} [\hat{f}(x)] - f(x))^2 \right]$$

Expanding this further:

$$= \mathbb{E} \left[ (\hat{f}(x) - \mathbb{E} [\hat{f}(x)])^2 \right] + \mathbb{E} \left[ (\mathbb{E} [\hat{f}(x)] - f(x))^2 \right] + 2\mathbb{E} \left[ (\hat{f}(x) - \mathbb{E} [\hat{f}(x)]) (\mathbb{E} [\hat{f}(x)] - f(x)) \right]$$

Since  $\mathbb{E} [\epsilon A] = \mathbb{E} [\epsilon] \mathbb{E} [A]$ ,  $A$  and  $\epsilon$  are independent and  $\mathbb{E} [\epsilon] = 0$  we have  $\mathbb{E} [\epsilon A] = 0$  thus:

$$\mathbb{E} [\mathbb{E} [\hat{f}(x)] - \hat{f}(x)] = \mathbb{E} [\hat{f}(x)] - \mathbb{E} [\hat{f}(x)] = 0$$



# Bias-Variance Decomposition Proof (cont.)

Thus, the expected squared error becomes:

$$\mathbb{E}_{data} \left[ (\hat{f}(x_n) - y)^2 \right] = \text{Variance} + \text{Bias}^2 + \sigma^2$$

where:

- **Variance** is  $\mathbb{E} \left[ (\hat{f}(x) - \mathbb{E} [\hat{f}(x)])^2 \right]$
- **Bias** is  $\mathbb{E} \left[ [\hat{f}(x)] - f(x) \right]$
- **Noise** is  $\sigma^2$

# High Bias in Simple Models

**Explanation:** Simple models, such as linear regression, often underfit

$$h_w(x) = w_0 + w_1 x$$

- Bias remains large even with infinite data

$$\text{Bias}^2 \gg \text{Variance}$$

- Leads to large generalization error

# High Variance in Complex Models

**Explanation:** Complex models tend to overfit

$$h_w(x) = w_0 + w_1x + w_2x^2 + \cdots + w_mx^m$$

- Variance dominates when the model is too complex

Variance  $\gg$  Bias

- Fits noise, leading to high test error

# Bias-Variance Tradeoff

**Tradeoff:** Balancing between bias and variance is key for optimal performance

- Low complexity: High bias, low variance
- High complexity: Low bias, high variance

# Regularization

**Purpose:** Prevent overfitting by penalizing large weights

$$J_{\lambda}(w) = J(w) + \lambda R(w)$$

- Common regularizers: L1 and L2 norms
- $\lambda$  controls the balance between fit and simplicity


# Effect of Regularization Parameter $\lambda$

## Balancing Fit and Complexity:

$$J_{\lambda}(w) = J(w) + \lambda \sum_{j=1}^m w_j^2 = J(w) + \lambda \mathbf{w}^T \mathbf{w}$$

- Large  $\lambda$ : Forces smaller weights, reduces complexity, increases bias, decreases variance
- Small  $\lambda$ : Allows larger weights, increases complexity, reduces bias, increases variance

# Effect of Regularization parameter $\lambda$



$$J_{\lambda}(w) = \sum_{i=1}^n (t^{(n)} - f(\mathbf{x}^{(n)}; \mathbf{w}))^2 + \lambda \mathbf{w}^T \mathbf{w}$$


$$f(\mathbf{x}^{(n)}; \mathbf{w}) = w_0 + w_1 x + \dots + w_9 x^9$$

# Effect of regularization on weights

	$\ln \lambda = -\infty$	$\ln \lambda = -18$	$\ln \lambda = 0$
$w_0^*$	0.35	0.35	0.13
$w_1^*$	232.37	4.74	-0.05
$w_2^*$	-5321.83	-0.77	-0.06
$w_3^*$	48568.31	-31.97	-0.05
$w_4^*$	-231639.30	-3.89	-0.03
$w_5^*$	640042.26	55.28	-0.02
$w_6^*$	-1061800.52	41.32	-0.01
$w_7^*$	1042400.18	-45.95	-0.00
$w_8^*$	-557682.99	-91.53	0.00
$w_9^*$	125201.43	72.68	0.01

Table adapted from Machine Learning and Pattern Recognition, Bishop





- $\lambda$  controls the effective complexity of the model
- hence the degree of overfitting

- 1 Generalization
- 2 Probabilistic regression
- 3 References

# Introduction to Regression (Probabilistic Perspective)

- **Objective:** Model the relationship between input  $\mathbf{x}$  and output  $y$ .
- **Uncertainty:** Output  $y$  has an associated uncertainty modeled by a probability distribution.

- **Example:**

$$y = f(\mathbf{x}; \mathbf{w}) + \epsilon \quad , \quad \epsilon \sim \mathcal{N}(0, \sigma^2)$$

- The goal is to learn  $f(\mathbf{x}; \mathbf{w})$  to predict  $y$ .

# Curve Fitting with Noise

- In real-world scenarios, observed output  $y$  is noisy.
- **Model: True output plus noise**

$$y = f(\mathbf{x}; \mathbf{w}) + \epsilon$$

- Noise represents unknown or unmodeled factors.
- **Example:** Predicting house prices based on features with inherent unpredictability.

# Expected Value of Output

- Best Estimate: The conditional expectation of  $y$  given  $\mathbf{x}$ .

$$\mathbb{E}[y|\mathbf{x}] = f(\mathbf{x}; \mathbf{w})$$

- Goal: Learn a function  $f(\mathbf{x}; \mathbf{w})$  that represents the average behavior of the data.
- Key Point: The model captures the mean of the target variable given input  $\mathbf{x}$ .

# Maximum Likelihood Estimation (MLE)

- **MLE:** A method to estimate parameters that maximize the likelihood of the data.
- Given data  $\mathcal{D} = \{(\mathbf{x}_i, y_i)\}_{i=1}^n$ , MLE maximizes:

$$L(\mathcal{D}; \mathbf{w}, \sigma^2) = \prod_{i=1}^n p(y_i | \mathbf{x}_i, \mathbf{w}, \sigma^2)$$

- MLE finds parameters  $\mathbf{w}$  and  $\sigma^2$  that best explain the data.

## Maximum Likelihood Estimation (cont.)

- Instead of maximizing the likelihood, it is often easier to maximize the log-likelihood:

$$\log L(\mathcal{D}; \mathbf{w}, \sigma^2) = \sum_{i=1}^n \log p(y_i | \mathbf{x}_i, \mathbf{w}, \sigma^2)$$

- It is because  $\log f(x)$  preserves the behaviour of  $f(x)$ .
- It is also easier to find derivative on summation of terms.

# Univariate Linear Function Example

- Assuming Gaussian noise with parameters  $(0, \sigma^2)$ , probability of observing real output value  $y$  is:

$$p(y|\mathbf{x}, \mathbf{w}, \sigma^2) = \frac{1}{\sqrt{2\pi\sigma^2}} \exp\left(-\frac{(y - f(\mathbf{x}; \mathbf{w}))^2}{2\sigma^2}\right)$$

- For a simple linear model  $f(\mathbf{x}; \mathbf{w}) = w_0 + w_1 x$  we have:

$$p(y|x, \mathbf{w}, \sigma^2) = \frac{1}{\sqrt{2\pi\sigma^2}} \exp\left(-\frac{(y - w_0 - w_1 x)^2}{2\sigma^2}\right)$$

- Key Observation:** Points far from the fitted line will have a low likelihood value.



## Log-Likelihood and Sum of Squares

- Using log-likelihood we have:

$$\log L(\mathcal{D}; \mathbf{w}, \sigma^2) = -n \log \sigma - \frac{n}{2} \log(2\pi) - \frac{1}{2\sigma^2} \sum_{i=1}^n (y^{(i)} - f(\mathbf{x}^{(i)}; \mathbf{w}))^2$$

- Since the objective of MLE is to optimize with regards to random variables, we can rule out the constants:

$$\log L(\mathcal{D}; \mathbf{w}, \sigma^2) \sim - \sum_{i=1}^n (y^{(i)} - f(\mathbf{x}^{(i)}; \mathbf{w}))^2$$

- Equivalence:** Maximizing the log-likelihood is equivalent to minimizing the Sum of Squared Errors (SSE):

$$J(\mathbf{w}) = \sum_{i=1}^n (y^{(i)} - f(\mathbf{x}^{(i)}; \mathbf{w}))^2$$

# Estimating $\sigma^2$

- The maximum likelihood estimate of the noise variance  $\sigma^2$ :

$$\hat{\sigma}^2 = \frac{1}{n} \sum_{i=1}^n \left( y^{(i)} - f(\mathbf{x}^{(i)}; \hat{\mathbf{w}}) \right)^2$$

- Interpretation: Mean squared error of the predictions.
- Note:  $\sigma^2$  reflects the noise level in the observations.

- 1 Generalization
- 2 Probabilistic regression
- 3 References

# Contributions

- **These slides are authored by:**
  - Arshia Gharooni
  - Mahan Bayhaghi

- [1] C. M., *Pattern Recognition and Machine Learning*.  
Information Science and Statistics, New York, NY: Springer, 1 ed., Aug. 2006.
- [2] M. Soleymani Baghshah, “Machine learning.” Lecture slides.
- [3] A. Ng and T. Ma, *CS229 Lecture Notes*.
- [4] T. Mitchell, *Machine Learning*.  
McGraw-Hill series in computer science, New York, NY: McGraw-Hill Professional,  
Mar. 1997.
- [5] Y. S. Abu-Mostafa, M. Magdon-Ismail, and H.-T. Lin, *Learning From Data: A Short Course*.  
New York, NY: AMLBook, 2012.
- [6] S. Goel, H. Bansal, S. Bhatia, R. A. Rossi, V. Vinay, and A. Grover, “CyCLIP: Cyclic Contrastive Language-Image Pretraining,” *ArXiv*, vol. abs/2205.14459, May 2022.